

CS220 - Computer System II
Lab 6

Due: 10/05/2017, 11:59pm

1 Introduction

In this lab, you will play with **pointers** to modify contents of the stack to corrupt data and control flow.

2 Data Corruption: The Problem

In Prog.c given below, you are to implement **bad** function such that the print statement prints the name of the student as "Quick brown fox jumped over the lazy dog" and age as 1000.

```
1  /* Prog.c */
   #include <stdio.h>
3  void bad()
   {
5     /* Your implementation goes here */
   }
7
   int main() {
9     struct {
        char *name;
11        int age;
    } student = {.name = "John", .age = 22 };
13    bad();
    printf("student.name = %s\nstudent.age = %d\n", student.name, student.age);
15    return 0;
   }
```

RULES:

1. You are not allowed to change main. All your changes must be confined in the bad function.
2. bad function is not allowed to accept any input from user. That is, you are not allowed to manually pass the address of **student** to **bad** function.

3 The Solution

3.1 Strategy

We know that the local variables are stored on the stack. If we can **locate the student structure on the stack**, then, we can modify the “name” field to point to any string. Similarly, we can modify the age field to a value 1000.

3.2 Step 1: Creating the new name

First, let us create a new string in `bad` with the new value. We will also create a pointer called `temp` that we will use. We use type `unsigned char *` for `temp` to allow us to perform pointer arithmetic easily. We initialize it to the address of `new_name`.

```
void bad() {  
2   char *new_name = "Quick brown fox jumped over the lazy dog";  
   unsigned char *temp;  
4   temp = (unsigned char *) (&new_name);  
   ...  
6 }
```

Note that although `name` is local to `bad`, because it is a string literal, **it is created in the data section, so its scope is not limited to the scope of `bad` function.**

3.3 Step 2: Locating address of student and new_name

If the student structure is in `main`’s stack frame, the `new_name` variable is located in `bad` function’s stack frame, which is **on top of `main`’s stack frame**. So, **`&student` is at positive offset from `&new_name`**. The offset is positive because the stack grows from higher to lower address. That is, the top of the stack is at a lower address than bottom of the stack. Let us use `gdb` to find the some addresses. We are interested in address of `student.name`, `student.age` and `new_name`. We are also interested in finding the location where our string is stored, which is nothing but the value of `new_name`.

```
$ gcc -g prog.c -o prog -O0 /* -O0 will turn off optimizations */  
2 $ gdb ./prog  
(gdb) b main
```

```

4 (gdb) r
(gdb) p &student.name
6 (gdb) p &student.age
(gdb) b bad
8 (gdb) c
(gdb) p &new_name
10 (gdb) n /* This will initialize new_name with the string */

```

Record the addresses of `student.name`, `student.age`, and `new_name`. We also need the location where the new string is stored, but we don't need to separately record it.

3.4 Step 3: Computing the offset

Compute the offset of `student.name` from `new_name`, which is nothing but `&student.name - &new_name`. Record it as `Offset_name`. Similarly, record `Offset_age` as the offset from `&student.age` and `&new_name`.

3.5 Step 4: Modifying the values

The final step is to make `student.name` to point to the new string and change the value at `student.age` to 1000.

```

void bad()
2 {
  char *new_name = "Quick brown fox jumped over the lazy dog";
4  unsigned char *temp;
  temp = (unsigned char *) (&new_name);
6
  /* We are telling the compiler: Treat as if temp+offset_name contains an
    address to a string and set the address of the new_string as the contents
    of the location */
8  *((char **) (temp + Offset_name)) = new_name;

  /* Treat as if temp+offset_age is a pointer to an integer and store the
10  value 1000 in the location pointed to by that integer. */
  *((int *) (temp + Offset_age)) = 1000;
12 }

```

4 Altering control-flow

Going one step further, because the return address is located on the stack, the problem to ask is: can we `modify the return address` to return to some other location. The answer is yes. In the above program, the return from `bad` function returns control to the instruction/statement immediately after the call to `bad` function. Now, let us change that such that the print statement in main that prints the students name and age is skipped!

4.1 Strategy

Just like how we altered the name and age fields, we first want to find the location where the return address is stored on stack. Then, given the address where the return statement in main is, we want to change the return address from bad to point to the return statement in main so that the print statement in main is skipped.

4.2 Step 1: Finding the address that bad returns to

Use gdb and find the address `bad` *should* return to.

```
$ gdb ./prog
2 (gdb) b main
(gdb) r
4 (gdb) disas main
```

Locate the call instruction to bad. It looks like this:

```
addr <+offset from start of function>: callq <bad>
```

Note that `address of instruction immediately after call to bad`. Record it as `ret_orig`.

4.3 Step 2: Find the address print statement in main returns to

Next, we identify the new value of the return address. That is, we want the address that follows after the call to `printf`. Using similar steps as above, locate the address in main that follows call to `printf`. Call it `ret_new`. Compute the distance/offset from `ret_orig` and `ret_new`. Call it `ret_offset`.

4.4 Step 3: Compute offset from `new_name` to return address on stack

Use `gdb` to intercept control flow in the `bad` function and print the contents of the stack. Locate `ret_orig`. Record the address where `ret_orig` is stored. This is the address where return address is located.

```
$ gdb ./prog
2 (gdb) b bad
(gdb) r
4 (gdb) x/16gx $rsp
```

Here, `rsp` is the address of the top of the stack. We are printing 16 64bit values in hex starting from the top of the stack. This prints the first 16 64bit values on the stack. Search for `ret_orig` and record the address where it is located. Stack addresses typically start with `0x7fff`. Call it `addr_ret_orig`. Next, given the address of `new_name`, compute the distance to `addr_ret_orig`. Call it `stack_offset_ret_addr`.

4.5 Step 4: Finally modify the return address to point to `ret_new`

```
void bad() {
2 ...
temp = (unsigned char *) (&new_name);
4 *((unsigned long *) (temp + stack_offset_ret_addr)) += ret_offset;
}
```

4.6 Step 5: Testing

Finally, compile and run the program. If all went well, you should see main return without printing student's name and age.