

Virtual Memory

Aravind Prakash

Modified version of slides by:
Dr. Karthic Palaniappan

Agenda

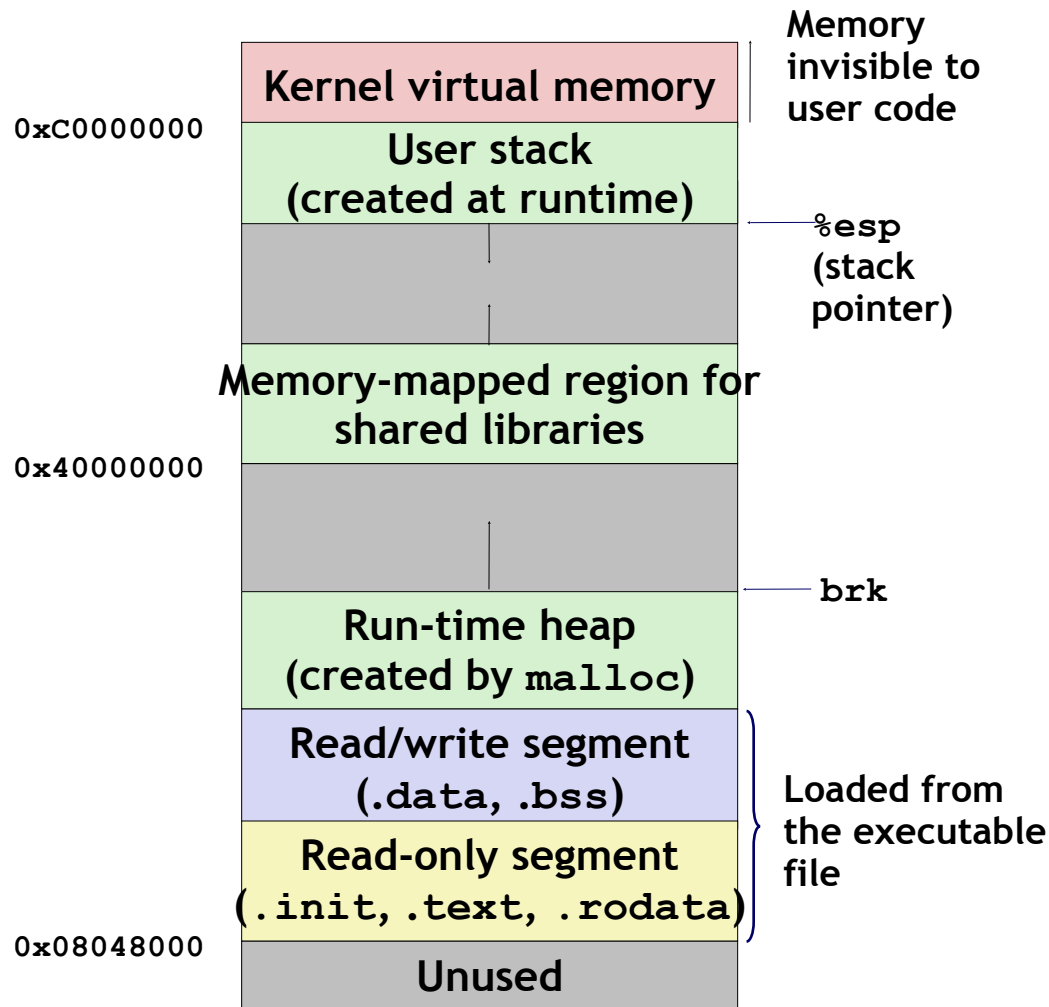
■ Virtual Memory Concepts

■ Address Translation

- Basic
- TLB
- Multilevel

Virtual Memory Concepts

- We've been viewing memory as a linear array.
- But wait! If you're running 5 processes with stacks at 0xC0000000, don't their addresses conflict?
- Nope! Each process has its own address space.
- How???



Virtual memory concepts

- We define a mapping from the **virtual** address used by the process to the actual **physical** address of the data in memory.

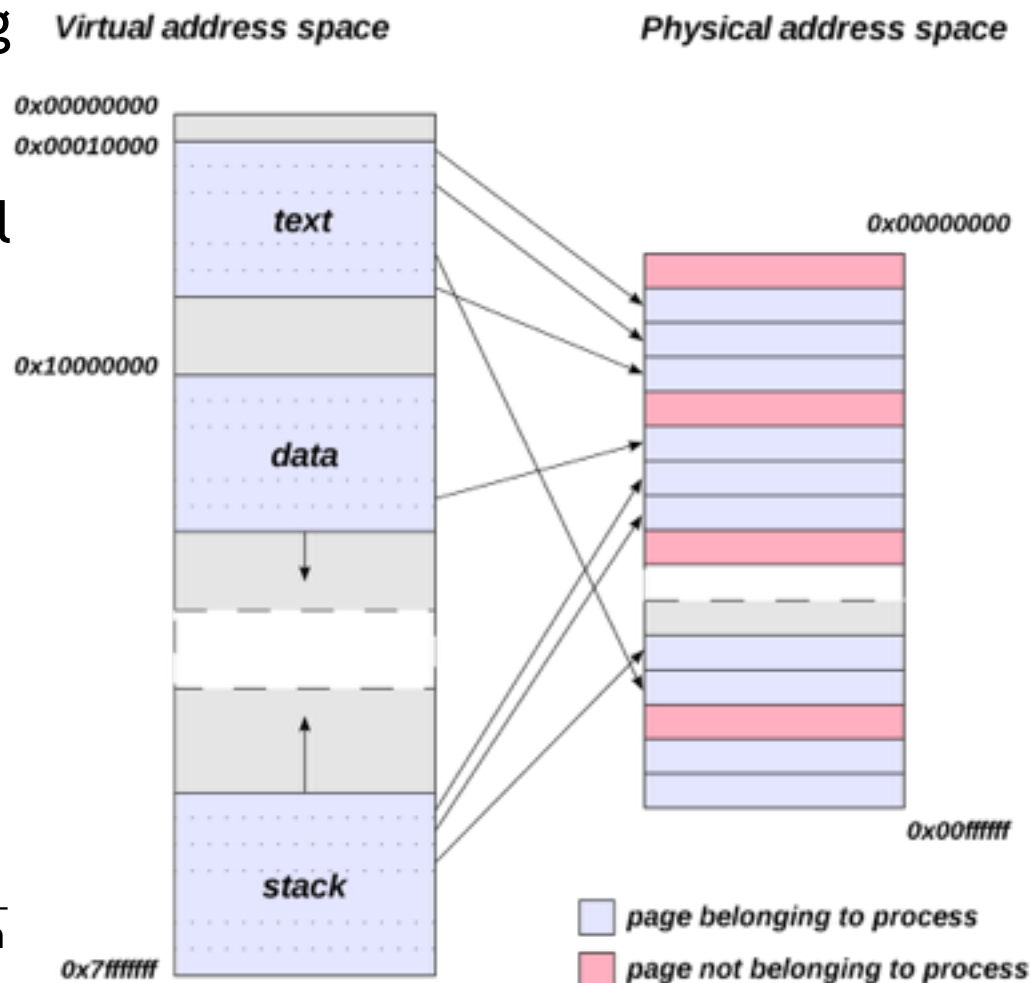
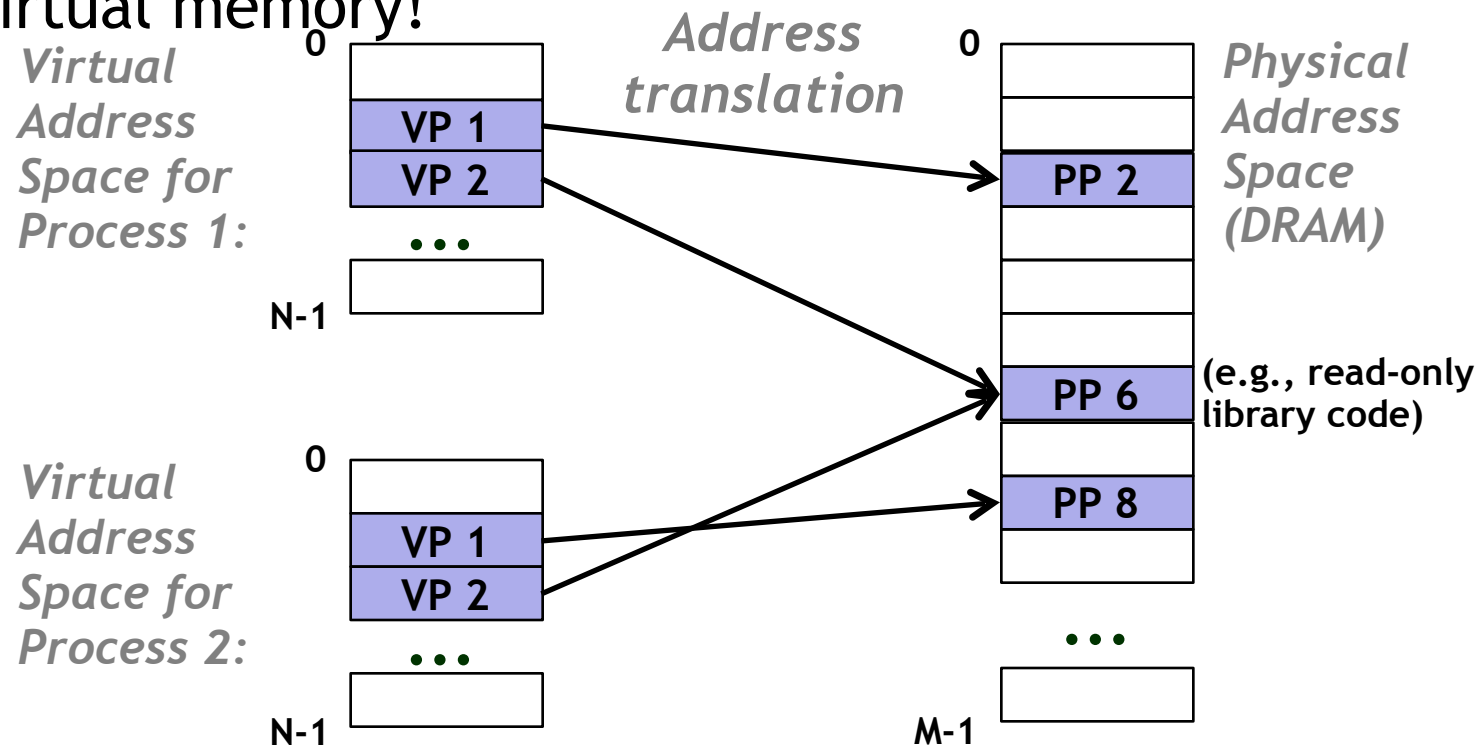


Image: http://en.wikipedia.org/wiki/File:Virtual_address_space_and_physical_address_space_relationship.svg

Virtual memory concepts

This explains why two different processes can use the same address. It also lets them share data *and* protects their data from illegal accesses. Hooray for virtual memory!



Virtual memory concepts

■ Page table

- Lets us look up the physical address corresponding to any virtual address. (Array of physical addresses, indexed by virtual address.)

■ TLB (Translation Lookaside Buffer)

- A special tiny cache just for page table entries.
- Speeds up translation.

■ Multi-level page tables

- The address space is often sparse.
- Use page directory to map large chunks of memory to a page table.
- Mark large unmapped regions as non-present in page directory instead of storing page tables full of invalid entries.

Agenda

■ Virtual Memory Concepts

■ Address Translation

- Basic
- TLB
- Multilevel

VM Address Translation

■ Virtual Address Space

- $V = \{0, 1, \dots, N-1\}$
- There are N possible virtual addresses.
- Virtual addresses are n bits long; $2^n = N$.

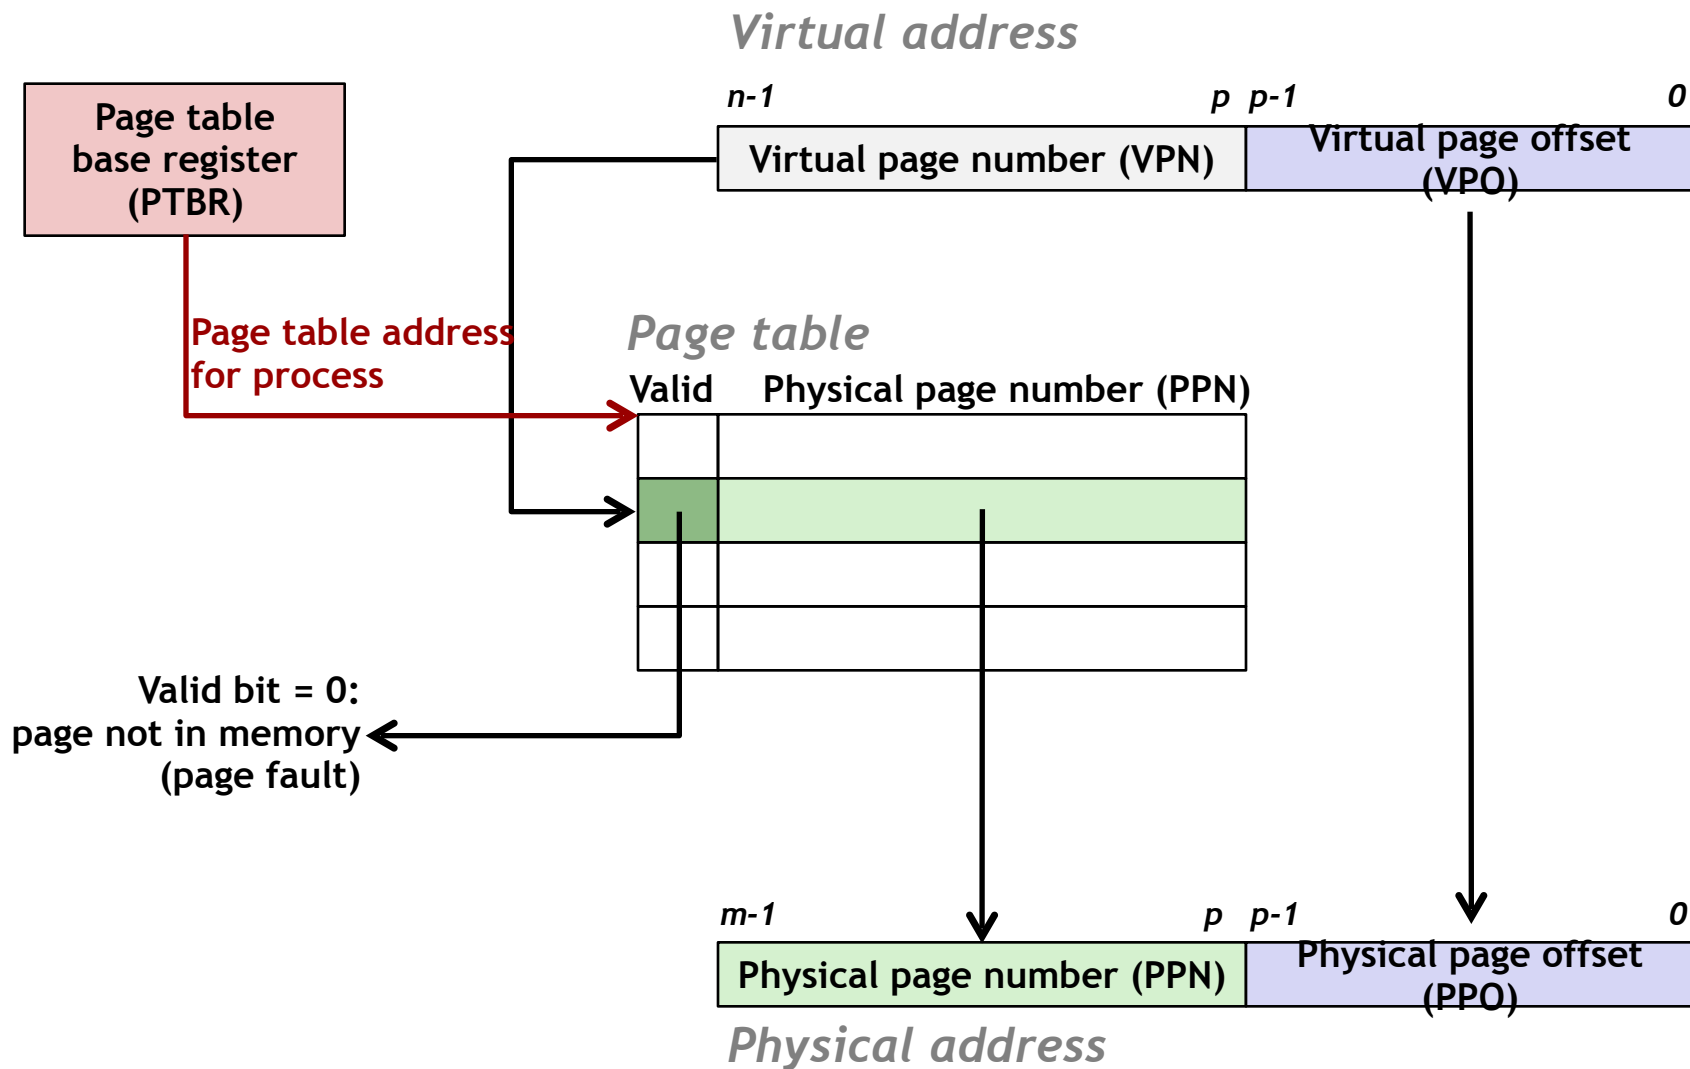
■ Physical Address Space

- $P = \{0, 1, \dots, M-1\}$
- There are M possible physical addresses.
- Virtual addresses are m bits long; $2^m = M$.

■ Memory is grouped into “pages.”

- Page size is P bytes.
- The address offset is p bytes; $2^p = P$.
- Since the virtual offset (VPO) and physical offset (PPO) are the same, the offset doesn't need to be translated.

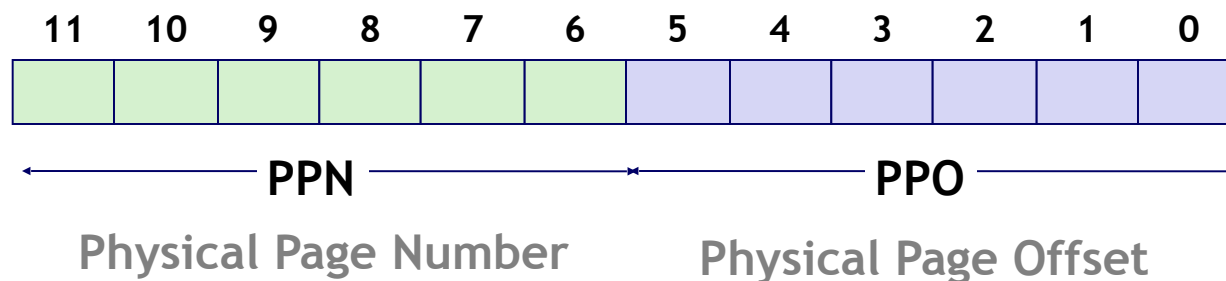
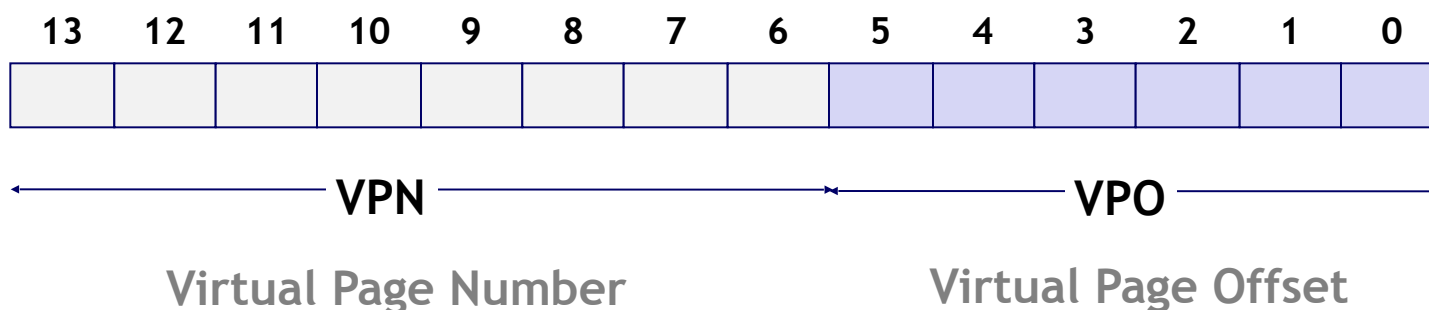
VM Address Translation



VM Address Translation

■ Addressing


- 14-bit virtual addresses
- 12-bit physical address
- Page size = 64 bytes



Example 1: Address Translation

■ Pages are 64 bytes. How many bits is the offset?

■ Find 0x03D4.




■ VPN: _____

■ PPN: _____

■ Physical address:

VPN	PPN	Valid
00	28	1
01	-	0
02	33	1
03	02	1
04	-	0
05	16	1
06	-	0
07	-	0


VPN	PPN	Valid
08	13	1
09	17	1
0A	09	1
0B	-	0
0C	-	0
0D	2D	1
0E	11	1
0F	0D	1



Example 1: Address Translation

■ Pages are 64 bytes. How many bits is the offset? $\log_2 64 = 6$

■ Find 0x03D4.




■ VPN: 0x0F

■ PPN: 0x0D

■ Physical address:
0x0354

VPN	PPN	Valid
00	28	1
01	-	0
02	33	1
03	02	1
04	-	0
05	16	1
06	-	0
07	-	0

VPN	PPN	Valid
08	13	1
09	17	1
0A	09	1
0B	-	0
0C	-	0
0D	2D	1
0E	11	1
0F	0D	1



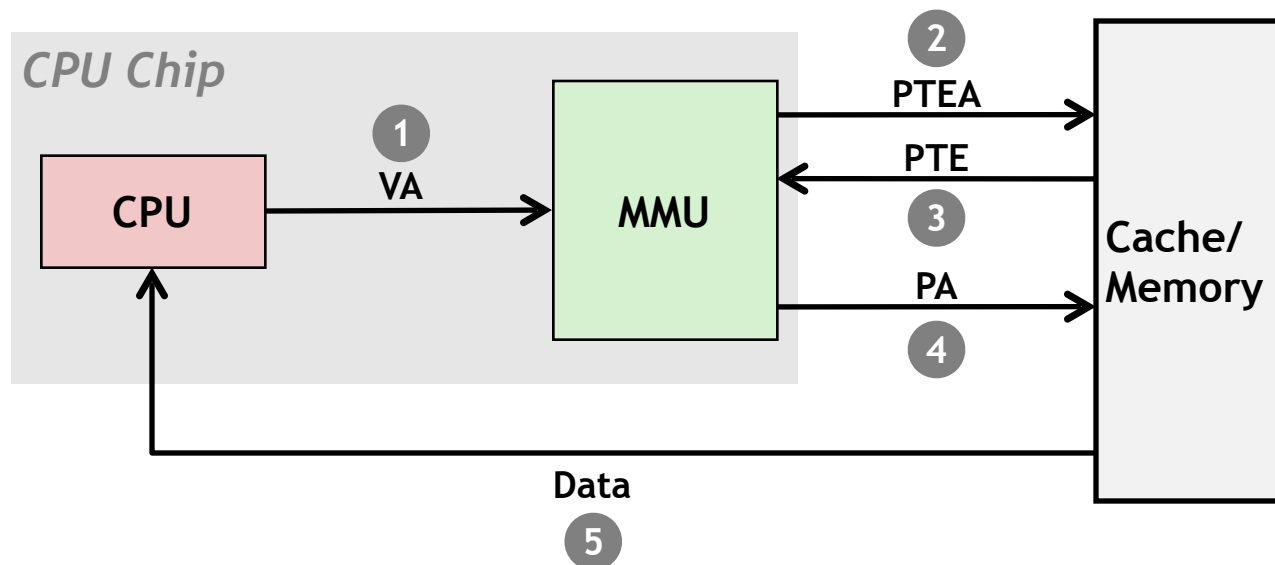
Agenda

■ Virtual Memory Concepts

■ Address Translation

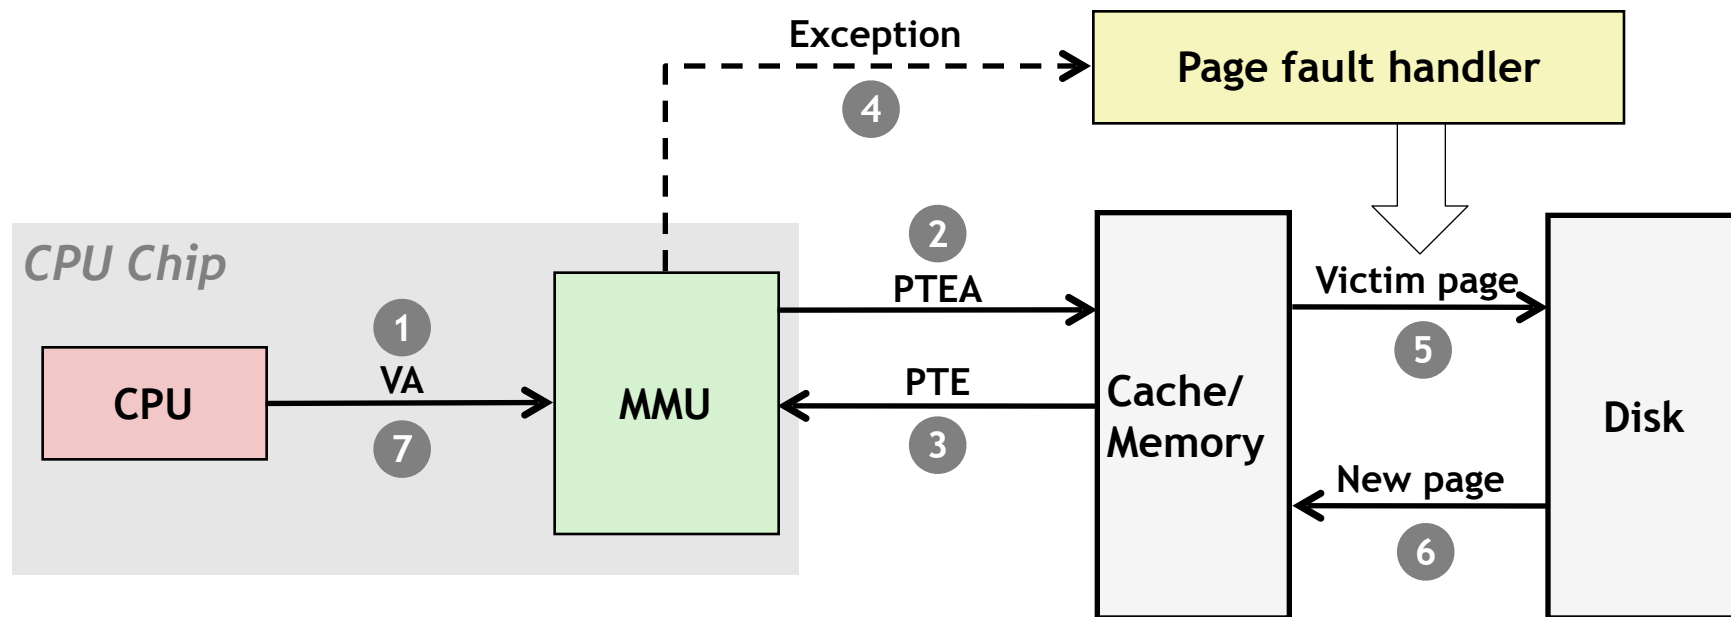
- Basic
- TLB
- Multilevel

Address Translation: Page Hit



- 1) Processor sends virtual address to MMU
- 2-3) MMU fetches PTE from page table in memory
- 4) MMU sends physical address to cache/memory
- 5) Cache/memory sends data word to processor

Address Translation: Page Fault



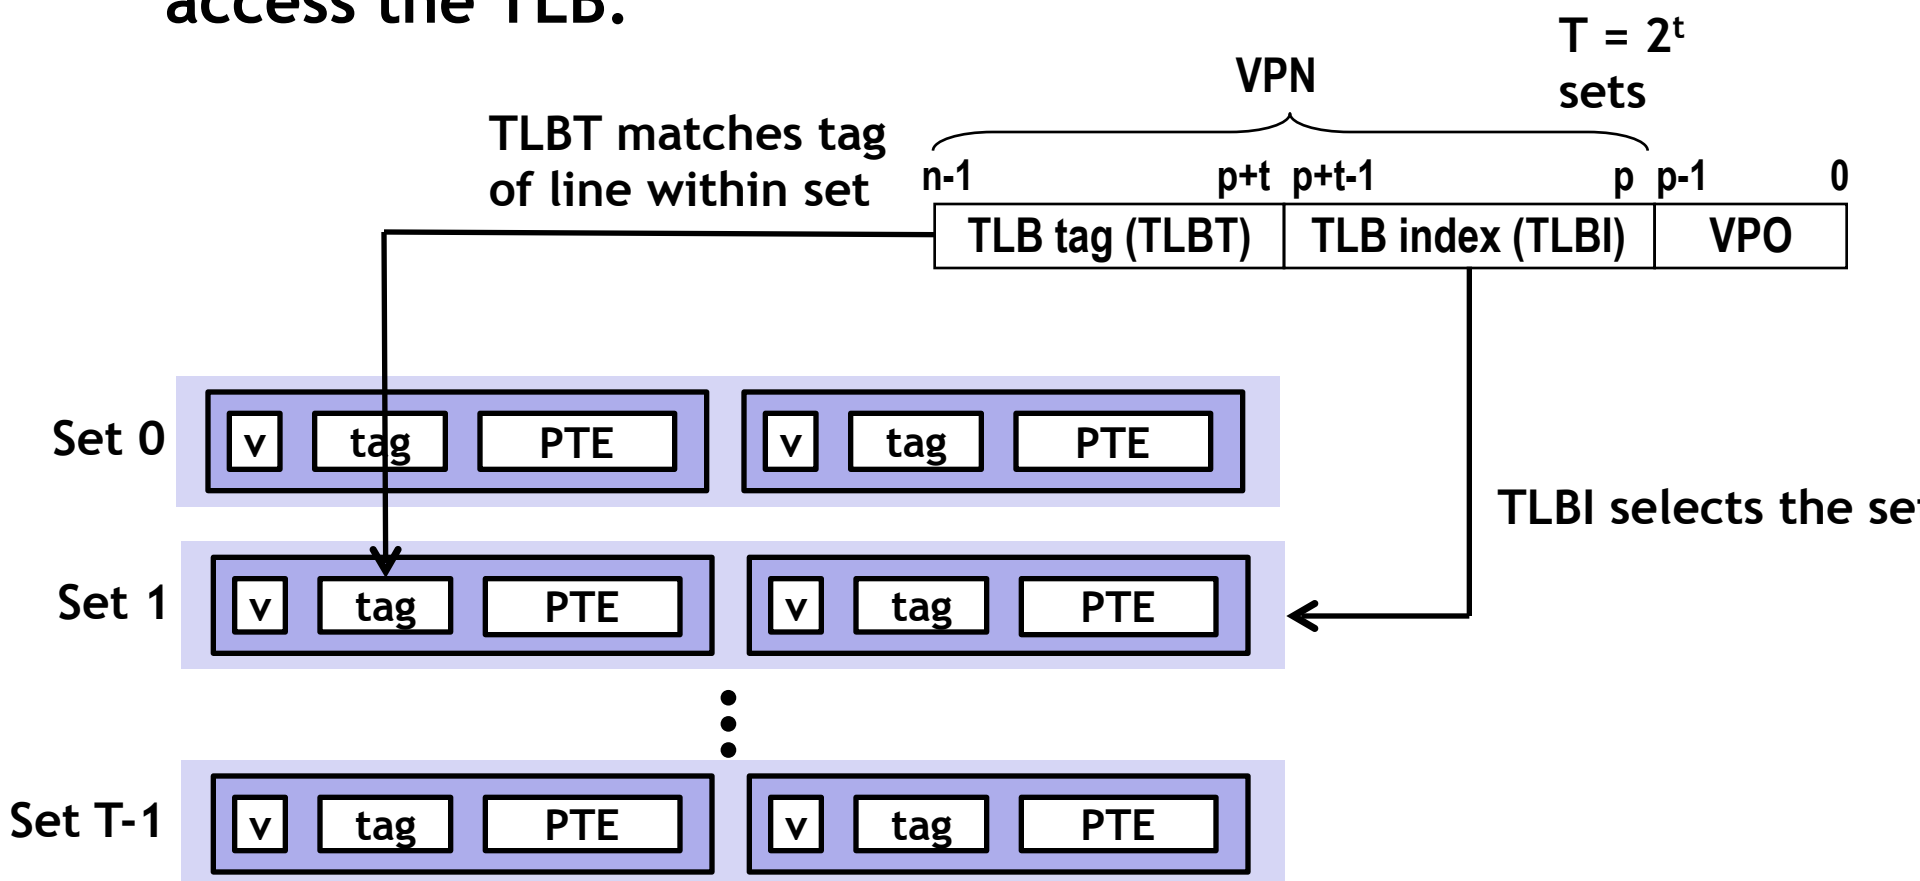
- 1) Processor sends virtual address to MMU
- 2-3) MMU fetches PTE from page table in memory
- 4) Valid bit is zero, so MMU triggers page fault exception
- 5) Handler identifies victim (and, if dirty, pages it out to disk)
- 6) Handler pages in new page and updates PTE in memory
- 7) Handler returns to original process, restarting faulting instruction

Speeding up Translation with a TLB

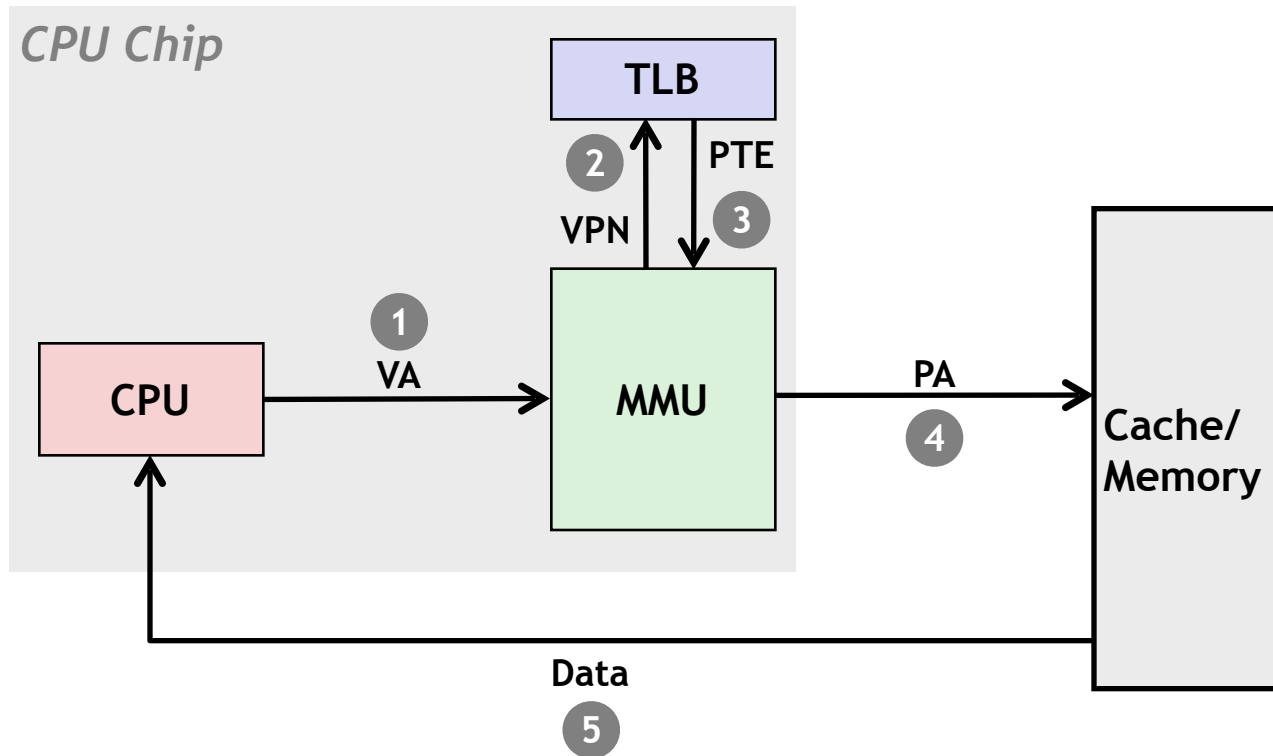
- Page table entries (PTEs) are cached in L1 like any other memory word
 - PTEs may be evicted by other data references
 - PTE hit still requires a small L1 delay
- Solution: *Translation Lookaside Buffer* (TLB)
 - Small set-associative hardware cache in MMU
 - Maps virtual page numbers to physical page numbers
 - Contains complete page table entries for small number of pages

Accessing the TLB

- MMU uses the VPN portion of the virtual address to access the TLB:

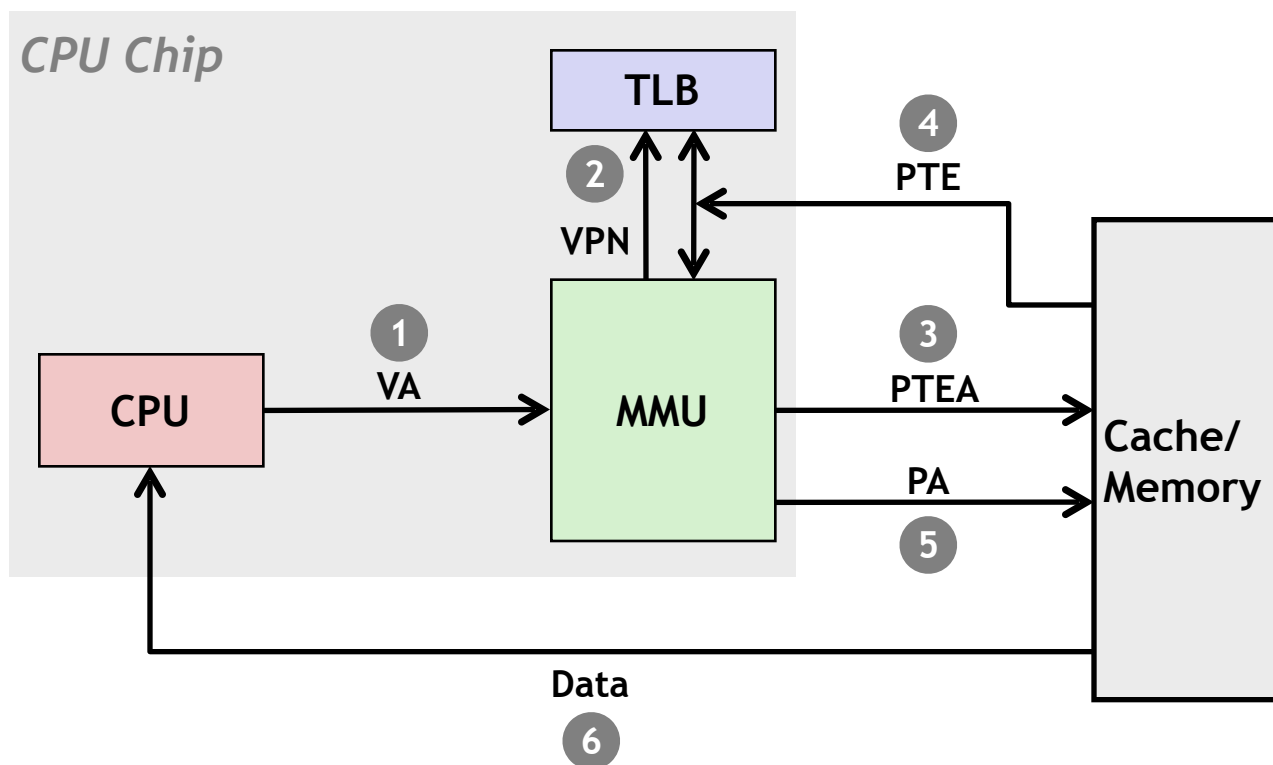


TLB Hit



A TLB hit eliminates a memory access

TLB Miss



A TLB miss incurs an additional memory access (the PTE)

Fortunately, TLB misses are rare. Why?

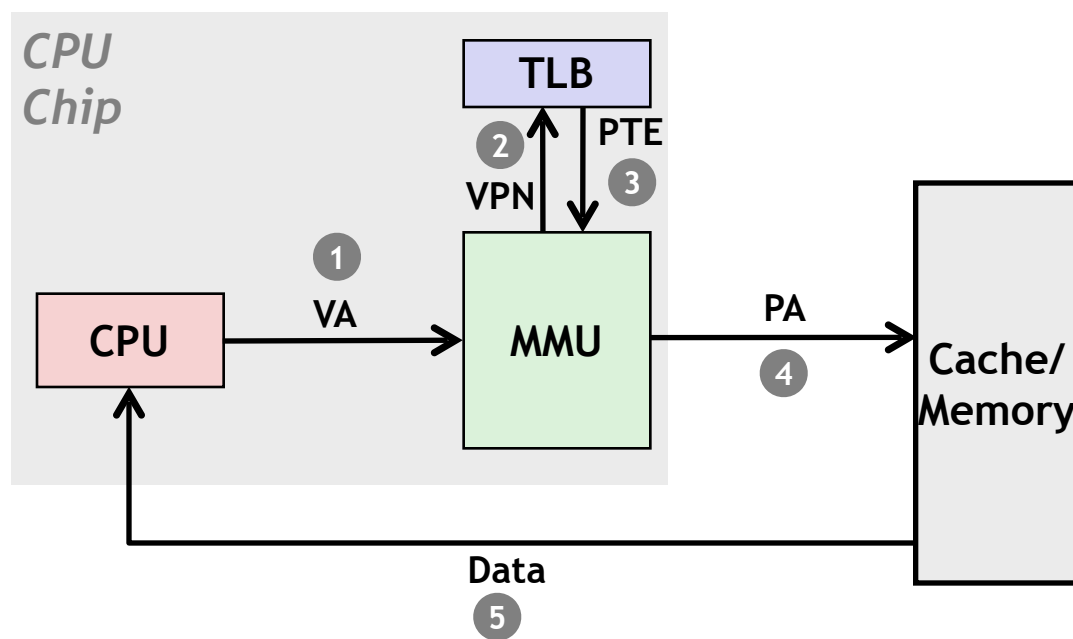
VM Address Translation with TLB

■ That's nice and simple, but it doubles memory usage.

- One memory access to look in the page table.
- One memory access of the actual memory we're looking for.

■ **Solution:**

- Cache the most frequently used page table entries in the TLB.
- To look up a virtual address in the TLB, split up the VPN (**not the whole virtual address!**) into a TLB index and a TLB tag.



Example 2: Address Translation with TLB

1 MB of virtual memory

4 KB page size

256 KB of physical memory

TLB: 8 entries, 2-way set associative

- How many bits are needed to represent the virtual address space?
- How many bits are needed to represent the physical address space?
- How many bits are needed to represent the offset?

- How many bits are needed to represent VPN?
- How many bits are in the TLB index?
- How many bits are in the TLB tag?

Example 2: Address Translation with TLB

1 MB of virtual memory

4 KB page size

256 KB of physical memory

TLB: 8 entries, 2-way set associative

- How many bits are needed to represent the virtual address space? **20. (1 MB = 2^{20} bytes.)**
- How many bits are needed to represent the physical address space? **18. (256 KB = 2^{18} bytes.)**
- How many bits are needed to represent the offset?
12. (4 KB = 2^{12} bytes.)
- How many bits are needed to represent VPN? **8. (20-12.)**
- How many bits are in the TLB index? **2. (4 sets = 2^2 set bits.)**
- How many bits are in the TLB tag? **6. (8-2.)**

Example 2a: Address Translation with TLB


- Translate 0x14213, given the contents of TLB and the first 32 entries of the page table below. (All the numbers are in hexadecimal.)

TLB			
Index	Tag	PPN	Valid
0	05	13	1
	3F	15	1
1	10	0F	1
	0F	1E	0
2	1F	01	1
	11	1F	0
3	03	2B	1
	1D	23	0

Page Table					
VPN	PPN	Valid	VPN	PPN	Valid
00	17	1	10	26	0
01	28	1	11	17	0
02	14	1	12	0E	1
03	0B	0	13	10	1
04	26	0	14	13	1
05	13	0	15	1B	1
06	0F	1	16	31	1
07	10	1	17	12	0
08	1C	0	18	23	1
09	25	1	19	04	0
0A	31	0	1A	0C	1
0B	16	1	1B	2B	0
0C	01	0	1C	1E	0
0D	15	0	1D	3E	1
0E	0C	0	1E	27	1
0F	2B	1	1F	15	1

Example 2a: Address Translation with TLB

0x14213



VPN:
TLBI:
TLBT:

TLB			
Index	Tag	PPN	Valid
0	05	13	1
	3F	15	1
1	10	0F	1
	0F	1E	0
2	1F	01	1
	11	1F	0
3	03	2B	1
	1D	23	0


TLB Hit!

PPN:

Offset:

Physical address:

Example 2a: Address Translation with TLB



VPN: 0x14
 TLBI: 0x00
 TLBT: 0x05

TLB			
Index	Tag	PPN	Valid
0	05	13	1
	3F	15	1
1	10	0F	1
	0F	1E	0
2	1F	01	1
	11	1F	0
3	03	2B	1
	1D	23	0

TLB Hit!
 PPN: 0x13
 Offset: 0x213

Physical address:
0x13213

Example 2b: Address Translation with TLB


- Translate 0x1F213, given the contents of TLB and the first 32 entries of the page table below.

TLB			
Index	Tag	PPN	Valid
0	05	13	1
	3F	15	1
1	10	0F	1
	0F	1E	0
2	1F	01	1
	11	1F	0
3	03	2B	1
	1D	23	0

Page Table					
VPN	PPN	Valid	VPN	PPN	Valid
00	17	1	10	26	0
01	28	1	11	17	0
02	14	1	12	0E	1
03	0B	0	13	10	1
04	26	0	14	13	1
05	13	0	15	1B	1
06	0F	1	16	31	1
07	10	1	17	12	0
08	1C	0	18	23	1
09	25	1	19	04	0
0A	31	0	1A	0C	1
0B	16	1	1B	2B	0
0C	01	0	1C	1E	0
0D	15	0	1D	3E	1
0E	0C	0	1E	27	1
0F	2B	1	1F	15	1

Example 2b: Address Translation with TLB


0x1F213



VPN:
TLBI:
TLBT:

TLB			
Index	Tag	PPN	Valid
0	05	13	1
	3F	15	1
1	10	0F	1
	0F	1E	0
2	1F	01	1
	11	1F	0
3	03	2B	1
	1D	23	0

Example 2b: Address Translation with TLB




VPN: 0x1F
 TLBI: 0x03
 TLBT: 0x07

TLB			
Index	Tag	PPN	Valid
0	05	13	1
	3F	15	1
1	10	0F	1
	0F	1E	0
2	1F	01	1
	11	1F	0
3	03	2B	1
	1D	23	0

TLB Miss!

Step 2: look it up in the page table. ☹

Example 2b: Address Translation with TLB




VPN: 0x1F
 TLBI: 0x03
 TLBT: 0x07

Page Table					
VPN	PPN	Valid	VPN	PPN	Valid
00	17	1	10	26	0
01	28	1	11	17	0
02	14	1	12	0E	1
03	0B	0	13	10	1
04	26	0	14	13	1
05	13	0	15	1B	1
06	0F	1	16	31	1
07	10	1	17	12	0
08	1C	0	18	23	1
09	25	1	19	04	0
0A	31	0	1A	0C	1
0B	16	1	1B	2B	0
0C	01	0	1C	1E	0
0D	15	0	1D	3E	1
0E	0C	0	1E	27	1
0F	2B	1	<u>1F</u>	<u>15</u>	<u>1</u>

Page Table Hit
 PPN:
 Offset:

Physical address:

Example 2b: Address Translation with TLB



VPN: 0x1F
TLBI: 0x03
TLBT: 0x07

Page Table					
VPN	PPN	Valid	VPN	PPN	Valid
00	17	1	10	26	0
01	28	1	11	17	0
02	14	1	12	0E	1
03	0B	0	13	10	1
04	26	0	14	13	1
05	13	0	15	1B	1
06	0F	1	16	31	1
07	10	1	17	12	0
08	1C	0	18	23	1
09	25	1	19	04	0
0A	31	0	1A	0C	1
0B	16	1	1B	2B	0
0C	01	0	1C	1E	0
0D	15	0	1D	3E	1
0E	0C	0	1E	27	1
0F	2B	1	<u>1F</u>	<u>15</u>	<u>1</u>

Page Table Hit

PPN: 0x15

Offset: 0x213

Physical address:
0x15213

Agenda

■ Virtual Memory Concepts

■ Address Translation

- Basic
- TLB
- Multilevel

Address Translation in Real Life

- Multi level page tables, with the first level often called a “page directory”
- Use first part of the VPN to get to the right directory and then the next part to get the PPN
- K-level page table divides VPN into k parts

